Algorithms Notes

by Tyler Wright

github.com/Fluxanoia   fluxanoia.co.uk

These notes are not necessarily correct, consistent, representative of the course as it stands today or, rigorous. Any result of the above is not the author's fault.

These notes are marked as unsupported, they were supported up until June 2019.

1 Bounding

1.1 Racetrack Principle

For $f, g : \mathbb{N} \to \mathbb{N}$ functions, n, k in $\mathbb{N}$ we have that:

$$\left.\begin{array}{l} f(k) \geq g(k) \\ f'(n) \geq g'(n) \quad (\forall n \geq k) \end{array}\right\} \Rightarrow f(n) \geq g(n) \quad (\forall n \geq k)$$

If a function f is greater than another function g at a value k and has a greater gradient for all values after and including k, f is greater than g for all values after and including k.

1.2 Big O Notation

1.2.1 Definition of the big O notation

For $g : \mathbb{N} \to \mathbb{N}$ a function, $O(g)$ is a set of functions $f : \mathbb{N} \to \mathbb{N}$ such that each for f in $O(g)$:

$$\exists c \in \mathbb{R}_{>0}, n_0 \in \mathbb{N} \text{ such that } \forall n \in N,$$
$$(n \geq n_0) \Rightarrow (0 \leq f(n) \leq cg(n)).$$

1.2.2 The big O notation under multiplication

For $f_1, f_2, g_1, g_2 : \mathbb{N} \to \mathbb{N}$ functions where:

- $f_1 \in O(g_1)$

- $f_2 \in O(g_2)$,

we have that:

- $f_1 + f_2$ is in $O(g_1 + g_2)$

- $f_1 \cdot f_2$ is in $O(g_1 \cdot g_2)$.

1.2.3 Closure of the big O notation

For $g : \mathbb{N} \to \mathbb{N}$ a function, $O(g)$ is closed under addition (this follows from the above).

1.2.4 Polynomials and the big O notation

For $p : \mathbb{N} \to \mathbb{N}$ a polynomial of degree k, p is in $O(n^k)$.

1.3   Θ Notation

1.3.1   Definition of the Θ notation

For $g : \mathbb{N} \to \mathbb{N}$ a function, $\Theta(g)$ is a set of functions $f : \mathbb{N} \to \mathbb{N}$ such that each for f in $\Theta(g)$:

$$\exists\, c_1, c_2 \in \mathbb{R}_{>0}, n_0 \in \mathbb{N} \text{ such that } \forall n \in N,$$
$$(n \geq n_0) \Rightarrow (0 \leq c_1 g(n) \leq f(n) \leq c_2 g(n)).$$

f is sandwiched by multiples of g.

1.3.2   Equivalency of the Θ notation

For $f, g : \mathbb{N} \to \mathbb{N}$ functions:

$$f \in \Theta(g) \iff g \in \Theta(f).$$

1.3.3   Θ and O notation

For $f, g : \mathbb{N} \to \mathbb{N}$ functions:

$$f \in \Theta(g) \iff f \in O(g) \text{ and } g \in O(f).$$

1.3.4   Definition of the Ω notation

For $g : \mathbb{N} \to \mathbb{N}$ a function, $\Omega(g)$ is a set of functions $f : \mathbb{N} \to \mathbb{N}$ such that each for f in $\Omega(g)$:

$$\exists\, c \in \mathbb{R}_{>0}, n_0 \in \mathbb{N} \text{ such that } \forall n \in N,$$
$$(n \geq n_0) \Rightarrow (0 \leq cg(n) \leq f(n)).$$

1.3.5   Equivalency of the Ω notation

For $f, g : \mathbb{N} \to \mathbb{N}$ functions:

$$f \in \Omega(g) \iff g \in O(f).$$

2 Runtime

2.1 Best-case Runtime

Considering all the inputs for a given algorithm, the best-case runtime is the runtime of the input that the algorithm takes the least amount of time to process.

2.2 Worst-case Runtime

Considering all the inputs for a given algorithm, the worst-case runtime is the runtime of the input that the algorithm takes the most amount of time to process.

2.3 Average Runtime

Considering all the inputs for a given algorithm, the average runtime is the average of the runtimes of all the inputs.

3 Data Structures

3.1 Trees

3.1.1 Definition of a tree

A tree T of size n is defined as $T = (V, E)$ where:

$$V = \{v_1, \ldots, v_n\} \text{ is a set of nodes}$$
$$E = \{e_1, \ldots, e_{n-1}\} \text{ is a set of edges,}$$

with the properties that for i in $\{1, \ldots, n-1\}$, j, k in $\{1, \ldots, n\}$ with $j \neq k$ we have $e_i = \{v_j, v_k\}$, and for all i in $\{1, \ldots, n\}$, there exists j in $\{1, \ldots, n-1\}$ such that v_i is in e_j.

Basically, we have n nodes and $n - 1$ edges where each node has least one edge and the edges can't branch between identical nodes.

3.1.2 Rooted trees

A rooted tree is defined as $T = (v, V, E)$ where $T = (V, E)$ is a tree and v in V is the root of T.

3.1.3 Leaves and internal nodes

A leaf in a tree is a node with exactly one incident edge. If a node isn't a leaf, it's an internal node.

3.1.4 Other definitions

- The **parent** of a node is the closest node on the path from the node to the root (the root has no parent)

- The **children** of a node are all its neighbours barring its parent

- The **height** of a tree is the length of the longest path connecting the root with a leaf

- The **degree** of a node is the number edges incident on the node

- The **level** of a node is the length of the unique path from the root to it plus one.

3.1.5 k-ary trees

A k-ary tree is a rooted tree where each node has at most k children. A k-ary tree is:

- **Full** if all internal nodes have exactly k children

- **Complete** if all levels except the last are full.

- **Perfect** if all levels are full.

Complete and perfect k-ary trees have heights of $O(\log_k(n))$.

3.2 Priority Queues

3.2.1 Definition of a priority queue

Priority queues are data structures that allow for the creation of a data structure from an array of values and the extraction of said data structure's maximum.

3.2.2 A tree-oriented priority queue

An array could be interpreted by a priority queue as a complete binary tree with the indices from top to bottom, left to right. So, where applicable, for a node index i, the parent of the node has index $\lfloor i/2 \rfloor$ and the children have indices $2i$ and $2i+1$.

3.2.3 Heaps

If we have a tree-oriented priority queue and we add the condition (heap property) that the values of nodes must be greater than their children we get a tree with the maximum at the root. We call this tree a heap.

3.2.4 Producing a heap

Given a binary tree, we can transform it into a heap using a `heapify` function. This function takes a node and its children and ensures the maximum value of these nodes lies in the parent node.

<div align="center">

heapify

Input	A binary tree and an index such that the subtrees with the index's children as the roots are heaps
Output	A binary tree such that the subtree with the given index as the root is a heap
Runtime	$O(\log_2(n))$

</div>

We have that the runtime is $O(\log_2(n))$ as this is the order of the height of the tree (because when we swap, we call the function again on the swapped child).

Using this, we traverse through the internal nodes, performing `heapify`. If the function makes a change, we perform `heapify` on the child node that was swapped with. This produces a heap.

So, we can see that the heap building function has the following properties:

<div align="center">

Heap Building Function

Input	A binary tree
Output	A heap
Runtime	$O(n \log_2(n))$

</div>

*Note that this is **not** the Heap Sort function! This is purely a data structure construction function.*

4 Recurrences

4.1 Motivation

For an increasing function $T : \mathbb{N} \to \mathbb{N}$ that produces the worst-case runtime for an algorithm for an input n. If we have that T is defined in terms of itself:

$$T(1) = c_1$$
$$T(n) = c_2 T(f(n)) + c_3 g(n),$$

where c_1, c_2, c_3 are in $\mathbb{R}_{>0}$ and $f, g : \mathbb{N} \to \mathbb{N}$. We would like to bound this function.

4.2 Substitution

We can guess what bounds the function. Consider the following:

$$T(1) = c_1$$
$$T(n) = 2T(\lfloor n/2 \rfloor) + c_2 n.$$

Suppose we guess that $T(n) \leq C[n \log_2(n)]$, we will use induction to prove our claim (as we are looking for an asymptotic comparison we can take any n in $\mathbb{N}$ as our base case):

$$\text{Base case } (n = 2):$$

$$T(2) = 2T(1) + 2c_2 = 2(c_1 + c_2)$$
$$C[2 \log_2(2)] = 2C$$

$$\text{So, for } C \geq (c_1 + c_2), \ T(2) \leq C[2 \log_2(2)].$$

$$\text{Inductive step, suppose the claim holds for all } n < k:$$

$$\begin{aligned}
T(k) &= 2T(\lfloor k/2 \rfloor) + c_2(k) \\
&\leq 2(C[k/2] \log_2(k/2)) + c_2(k) \\
&= C[k \log_2(k/2)] + c_2(k) \\
&= Ck[\log_2(k) - 1] + c_2(k) \\
&= Ck \log_2(k) + k(c_2 - C)
\end{aligned}$$

$$\text{So, for } C \geq c_2, \text{ we have:}$$

$$T(k) \leq C[k \log_2(k)].$$

$$\text{So, for } C \geq (c_1 + c_2), \text{ the claim holds for all } n \geq 2 \text{ by induction.}$$

More educated guesses can be made by considering the branching, recursive calls of the algorithm.

5 Searching

5.1 Linear Search

5.1.1 Information on Linear Search

Input	An array of integers and an integer x both in $[0, n)$ for some n in $\mathbb{N}$
Output	1 if x is in the array, 0 otherwise
Best-case Runtime	$O(1)$
Average Runtime	$O(n)$
Worst-case Runtime	$O(n)$

5.1.2 Process of Linear Search

Iterate through the array comparing the input value with the current array value. If it's equal, return 1. If we reach the end of the array, return 0.

5.2 Binary Search

5.2.1 Information on Binary Search

Input	A sorted array of integers and an integer x both in $[0, n)$ for some n in $\mathbb{N}$
Output	1 if x is in the array, 0 otherwise
Best-case Runtime	$O(1)$
Average Runtime	$O(\log_2(n))$
Worst-case Runtime	$O(\log_2(n))$

5.2.2 Process of Binary Search

Look at the middle value of the array, if equal to the input value then return 1. If the value is greater than our input value, repeat the process with the lesser half of the array. Otherwise, repeat with the greater half of the array.

This works because the array is sorted.

6 Sorting

6.1 Properties of Sorting Algorithms

6.1.1 In place

A sorting algorithm is in place if at any moment at most $O(1)$ array elements are stored outside the array.

6.1.2 Stable

A sorting algorithm is stable if any pair of equal values appear in the same order in the sorted array (this may be important if this value is tied to some overarching data structure).

6.2 Lower Bound for Comparison-Based Sorting

When sorting comparatively, we can only sort an array length n at best in $O(n \log(n))$ time.

6.3 Insertion Sort

6.3.1 Information on Insertion Sort

Input	An array of integers length n in $\mathbb{N}$
Output	An ascending, sorted array
Best-case Runtime	$O(n)$
Average Runtime	$\Theta(n^2)$
Worst-case Runtime	$O(n^2)$
In place	✓
Stable	✓

6.3.2 Process of Insertion Sort

Iterate through the array A, when at position i, place $A[i]$ into the array at some index in $\{0, \ldots, i\}$ such that $A[0, i]$ is sorted.

6.4   Merge Sort

6.4.1   Information on Merge Sort

Input	An array of integers length n in $\mathbb{N}$
Output	An ascending, sorted array
Best-case Runtime	$O(n \log_2(n))$
Average Runtime	$O(n \log_2(n))$
Worst-case Runtime	$O(n \log_2(n))$
In place	$\times$
Stable	$\checkmark$

6.4.2   Process of Merge Sort

If the array size is less than 3, reorder the elements and return. Otherwise, split the array into two, perform merge sort on the two halves and combine them.

6.5   Heap Sort

6.5.1   Information on Heap Sort

Input	An array of integers length n in $\mathbb{N}$
Output	An ascending, sorted array
Runtime	$O(n \log_2(n))$
In place	$\checkmark$
Stable	$\times$

6.5.2   Process of Heap Sort

Produce a heap from the array and extract the maximum from it until it is empty. After each extraction, we swap the last element into the root and `heapify`.

6.6   Quick Sort

6.6.1   Information on Quick Sort

Input	An array of integers length n in $\mathbb{N}$
Output	An ascending, sorted array
Best-case Runtime	$O(n\log_2(n))$
Average Runtime	$O(n\log_2(n))$
Worst-case Runtime	$O(n^2)$
In place	✓
Stable	✗

(Based on a random pivot selection procedure)

6.6.2   Process of Quick Sort

Choose a pivot, move all values greater than the pivot into indices greater than the pivot's and move all values less than the pivot into indices less than the pivot.

This can be done by storing the smallest index i such that it's value has not been swapped then iterating through the array (excluding the pivot) and checking if the value is less or equal to the pivot. If so, move it to position i and increment i. Once at the end of the array, the pivot can be moved into position i.

Then, we perform quick sort on all values less than the pivot and on all values greater than the pivot.

6.7   Counting Sort

6.7.1   Information on Counting Sort

Input	An array of non-negative integers length n in $\mathbb{N}$
Output	An ascending, sorted array
Runtime	$O(n)$
In place	✗
Stable	✓

6.7.2   Process of Counting Sort

Create an array A with its length equal to the maximum of the input array. At each index, set the value equal to the amount of values less than the index in the input array. Then, we can fill our input array with the correct amount of each integer in increasing order.

6.8   Radix Sort

6.8.1   Information on Radix Sort

Input	An array of non-negative integers length n in $\mathbb{N}$
Output	An ascending, sorted array
Runtime	$O(n)$
In place	✓
Stable	✓

6.8.2   Process of Radix Sort

Using a stable sorting algorithm, sort the least to the most significant digits.

7   Dynamic Programming

7.1   Calculating Fibonacci Numbers

7.1.1   A simple algorithm

Input	An index
Output	The related Fibonacci number to the index
Runtime	$O(\phi^n)$

If the index is less than 2, we return the index. Otherwise, for indices $i \geq 2$, we sum this algorithm applied to $i - 1$ and $i - 2$.

7.1.2   An improved algorithm

Input	An index
Output	The related Fibonacci number to the index
Runtime	$O(n)$

We start at the beginning of the sequence (with 0 and 1). The next Fibonacci number can be calculated with the last two calculated Fibonacci terms. We repeat this until we get the Fibonacci number associated to our index.

7.2 Pole Cutting

7.2.1 Motivation

Supposed we have a pole of length n (in $\mathbb{N}$) that we can split into pieces of integral length and a function than produces a value for each pole length. How can we find out the best way to split the pole such that we maximise the value?

7.2.2 Information on Pole Cutting

Input	A pole length and a valuation function
Output	The best value possible
Runtime	$O(n^2)$

7.2.3 Process of Pole Cutting

For a valuation function f, the solution for a length n will be the maximum of the set of values:

$$[\text{optimal solution for } n-1] + f(1),$$
$$[\text{optimal solution for } n-2] + f(2),$$
$$\ldots$$
$$[\text{optimal solution for } 1] + f(n-1),$$
$$f(n)$$

As the optimal solution for $n = 1$ is trivial we can construct optimal solutions for greater values of n recursively until we have our desired optimal solution.

7.3 Matrix Chain Parenthesisation

7.3.1 Motivation

If we have a sequence of matrices to multiply, we would like to do this as fast as possible. By exploiting the associativity of matrix multiplication we reduce the number of necessary calculations.

7.3.2 Information on Matrix Chain Parenthesisation

Input	A sequence of matrices
Output	The optimal parenthesisation
Runtime	$O(n^3)$

7.3.3 Process of Matrix Chain Parenthesisation

For a sequence of matrices $\{A_1, A_2, \ldots, A_n\}$, we can split the sequence in two sequences of matrices $\{A_1, \ldots, A_k\}, \{A_{k+1}, \ldots, A_n\}$. Define:

$$A_{i \to j} = A_i A_{i+1} \cdots A_j$$
$$i, j \text{ in } \{1, \ldots, n\} \text{ with } i < j.$$

If we split the sequence at k in $\{0, \ldots, n\}$ ($A_0 = A_{n+1} = I$), the cost of this parenthesisation would be the cost of multiplying $A_{1 \to k}$ and $A_{k+1 \to n}$ plus the optimal parenthesisation cost of $A_{1 \to k}$ and $A_{k+1 \to n}$. So, if we produce optimal parenthesisation for each $A_{i \to j}$ with i, j in $\{1, \ldots, n\}$ and $i < j$ we can calculate the optimal solution for $A_{1 \to n}$.

8 Miscellaneous Algorithms

8.1 Peak Finding

8.1.1 Definition of a peak

A peak in an array of numbers is an index where its adjacent values are not greater than the value at said index.

8.1.2 Existence of a peak

Each array of numbers has a peak as the maximum is by definition a peak.

8.1.3 A simple algorithm

Input	An array of numbers
Output	An index
Runtime	$O(n)$

Check the end points of the array (to avoid over/underflow) and then iterate through the remaining indices, checking if each value is a peak. If a peak is found at any point in this process, return the index.

8.1.4 An improved algorithm

Input	An array of numbers
Output	An index
Runtime	$O(\log_2(n))$

If the array is length is 1 or 2, return the index of this maximum in the input array. Where the array is length $k > 2$, check if the value at $\lfloor k/2 \rfloor$ is a peak. If it is, return the index of this value in the input array. If it is not, it must have an adjacent value that is greater than it.

If the greater value is at index $\lfloor k/2 \rfloor + 1$, we perform this algorithm on the subarray $[\lfloor k/2 \rfloor + 1, k)$, if the greater value is at $\lfloor k/2 \rfloor - 1$, we perform this algorithm on the subarray $[0, \lfloor k/2 \rfloor)$.