

Types and Lambda Calculus Notes

by Tyler Wright

github.com/Fluxanoia

fluxanoia.co.uk

These notes are not necessarily correct, consistent, representative of the course as it stands today, or rigorous. Any result of the above is not the author's fault.

These notes are in progress.

Contents

1	The Axioms of Lambda Calculus	4
1.1	Axiom of Variables	4
1.2	Axiom of Application	4
1.3	Axiom of Abstraction	4
1.4	Subterms	4
1.5	Syntactical Conventions	5
2	Alpha	6
2.1	Free Variables	6
2.2	Substitution	6
2.3	Alpha Equivalence	6
2.4	The Variable Convention	6
3	Beta	7
3.1	Beta Reductions	7
3.2	Standard Combinators	7
3.3	Confluence of Beta	8
3.4	Beta Convertibility	8
4	Definability	9
4.1	Church Numerals	9
4.2	Lambda Definability	9
4.3	Basic Functions	9
4.4	Church Lists	9
5	Recursion	10
5.1	The Factorial	10
5.2	Fixed Points	10
5.3	The First Recursion Theorem	10
5.3.1	Solving Recursive Definitions	10
6	Induction	11
6.1	The Induction Metaprinciple	11
7	Types	12
7.1	Monotypes	12
7.2	Type Schemes	12
7.3	Free Type Variables	12
7.4	Type Substitution	12
7.4.1	Composition of Substitutions	12

7.5	Type Assignment	13
7.6	Type Environments	13
7.7	Type Judgement	13
7.8	The Type System	13
7.9	Typability	13
7.10	The Inversion Theorem	14
7.11	Properties of the System	14
7.12	Type Constraints	15
7.13	Constraint Generation	15
7.14	Unifier	15
	7.14.1 Most General Unifier	16
7.15	Constraint Solved Form	16
7.16	Robinson's Algorithm	16
7.17	Hindley-Milner Type Inference	16
7.18	Principle Type Scheme Theorem	16

1 The Axioms of Lambda Calculus

We suppose that we have a countably infinite set of variables \mathbb{V} (we usually refer to elements of this set as x, y, z , etc.), from this we define the alphabet of lambda calculus $\mathbb{V} + \{\lambda, ., (,)\}$.

The set of terms of lambda calculus Λ is defined inductively by the axioms. A string is a term if and only if there is a proof tree using the following axioms that concludes it is a term.

1.1 Axiom of Variables

For some x in \mathbb{V} , we have the axiom:

$$\overline{x \in \Lambda},$$

which says that each variable is a term.

1.2 Axiom of Application

We have the axiom:

$$\frac{M \in \Lambda \quad N \in \Lambda}{(MN) \in \Lambda},$$

which says that for some terms M and N , (MN) is also a term.

1.3 Axiom of Abstraction

For some x in \mathbb{V} , we have the axiom:

$$\frac{M \in \Lambda}{(\lambda x.M) \in \Lambda},$$

which says that for some term M , $(\lambda x.M)$ is also a term.

1.4 Subterms

Subterms of a term M are substrings of M that are themselves terms and not captured by a λ (directly preceded by).

1.5 Syntactical Conventions

Parentheses allow our lambda calculus to be unambiguous, but for the sake of simplicity, we will construct conventions that will allow us to retain unique meaning with less parentheses:

- Omit outermost parentheses,
- Terms associate to the left, (MNP) parses as $((MN)P)$,
- Bodies of abstractions end at parentheses, $(\lambda x.MN)$ parses as $(\lambda x.(MN))$,
- Group repeated abstractions, $(\lambda xy.M)$ parses as $(\lambda x.(\lambda y.M))$.

2 Alpha

2.1 Free Variables

We define the function $FV : \Lambda \rightarrow \mathcal{P}(\mathbb{V})$, which returns the set of variables contained within a term M that are not bound. We define it recursively on the structure of terms:

$$\begin{aligned} FV(x) &= \{x\}, \\ FV(MN) &= FV(M) \cup FV(N), \\ FV(\lambda x.M) &= FV(M) \setminus \{x\}. \end{aligned}$$

If a term has no free variables we say it is closed, and if a term has at least one free variable then we say it is open. The set of all closed terms is denoted by Λ^0 .

2.2 Substitution

We define 'capture-avoiding' substitution of a term M for a variable x recursively on the structure of terms:

$$\begin{aligned} y[M/x] &= y && \text{if } y \neq x, \\ y[M/x] &= M && \text{if } y = x, \\ (PQ)[M/x] &= P[M/x]Q[M/x], \\ (\lambda y.P)[M/x] &= \lambda y.P && \text{if } y = x, \\ (\lambda y.P)[M/x] &= \lambda y.P[M/x] && \text{if } y \neq x \text{ and } y \notin FV(M). \end{aligned}$$

On the final case, we stipulate that y cannot be a free variable of M because otherwise free variables in the substitution could be captured by the lambda.

2.3 Alpha Equivalence

Suppose we have a term $\lambda x.M$ and y in $\mathbb{V} \setminus FV(M)$, then substituting y for x is a change of bound variable name. If two terms can be made identical through changes of bound variable name, they are α -equivalent. The set of λ -terms is the set Λ under α -equivalence.

This equivalence is much more useful to us than string comparison, so for the remainder of the notes we will always be referring to λ -terms as terms.

2.4 The Variable Convention

For M_1, \dots, M_k terms occurring in the same scope, we assume each term has distinct bound variables. We can make this assumption as otherwise, we can use changes of bound variable names to make it so.

3 Beta

A term of the form $(\lambda x.M)N$ is called a β -redex and $M[N/x]$ is the contraction of the redex.

3.1 Beta Reductions

The one-step β -reduction relation, denoted by \rightarrow_β , is inductively defined by the redex rule:

$$\overline{(\lambda x.M)N \rightarrow_\beta M[N/x]},$$

the left and right application rules:

$$\frac{M \rightarrow_\beta M'}{MN \rightarrow_\beta M'N} \quad \frac{N \rightarrow_\beta N'}{MN \rightarrow_\beta MN'},$$

and the abstraction rule:

$$\frac{M \rightarrow_\beta N}{\lambda x.M \rightarrow_\beta \lambda x.N}.$$

A term M is said to be in β -normal form if there is no term N such that $M \rightarrow_\beta N$.

In general, β -reductions are sequences of consecutive one-step β -reductions:

$$M_0 \rightarrow_\beta M_1 \rightarrow_\beta \cdots \rightarrow_\beta M_k,$$

for some k in \mathbb{N}_0 . We say that M_0 β -reduces to M_k , denoted by $M_0 \rightarrow_\beta M_k$.

If $M \rightarrow_\beta N$, we say that N is a reduct of M , and is a proper reduct if $N \neq M$. If we can choose some β -normal N such that $M \rightarrow_\beta N$ then M is normalisable. If a term admits no infinite β -reductions then we say that it is strongly normalisable.

3.2 Standard Combinators

We have some interesting programs described below:

$$\begin{array}{ll} \mathbf{I} &= \lambda x.x \\ \mathbf{K} &= \lambda xy.x \\ \mathbf{S} &= \lambda xyz.xz(yz) \\ \omega &= \lambda x.xx \\ \Omega &= \omega\omega \\ \Theta &= (\lambda xy.y(xxy))(\lambda xy.y(xxy)) \end{array} \quad \begin{array}{l} \mathbf{I}M \rightarrow_\beta M \\ \mathbf{K}MN \rightarrow_\beta M \\ \mathbf{S}MNP \rightarrow_\beta MP(NP) \\ \omega M \rightarrow_\beta MM \\ \Omega \rightarrow_\beta \Omega \\ \Theta M \rightarrow_\beta M(\Theta M) \end{array}$$

3.3 Confluence of Beta

For a term M , if $M \rightarrow_{\beta} P$ and $M \rightarrow_{\beta} Q$ then there exists some term N such that $P \rightarrow_{\beta} N$ and $Q \rightarrow_{\beta} N$.

3.4 Beta Convertibility

For M, N terms, if M and N have a common reduct then we say that M and N are β -convertible, denoted by $M =_{\beta} N$.

4 Definability

4.1 Church Numerals

The Church numeral for the number n in \mathbb{N}_0 , denoted by $\ulcorner n \urcorner$, is:

$$\lambda f x. \underbrace{f(\cdots (f x) \cdots)}_{n \text{ times}}.$$

Each Church numeral is already in normal form.

4.2 Lambda Definability

A function $f : \mathbb{N} \times \cdots \times \mathbb{N} \rightarrow \mathbb{N}$ is λ -definable if there exists a λ -term F that satisfies:

$$F \ulcorner n_1 \urcorner \cdots \ulcorner n_k \urcorner =_{\beta} \ulcorner f(n_1, \dots, n_k) \urcorner.$$

4.3 Basic Functions

We have addition, predecessor, subtraction, and the zero conditional defined here:

Add	$= \lambda y z. \lambda f x. y f (z f x)$	$\mathbf{Add} \ulcorner m \urcorner \ulcorner n \urcorner =_{\beta} \ulcorner m + n \urcorner$
Pred	$= \lambda z. \lambda f x. z (\lambda g h. h (g f)) (\lambda u. x) (\lambda u. u)$	$\mathbf{Pred} \ulcorner 0 \urcorner =_{\beta} \ulcorner 0 \urcorner$ $\mathbf{Pred} \ulcorner n + 1 \urcorner =_{\beta} \ulcorner n \urcorner$
Sub	$= \lambda m n. n \mathbf{Pred} m$	$\mathbf{Sub} \ulcorner m \urcorner \ulcorner n \urcorner =_{\beta} \ulcorner 0 \urcorner$ if $m - n < 0$ $\mathbf{Sub} \ulcorner m \urcorner \ulcorner n \urcorner =_{\beta} \ulcorner m - n \urcorner$ otherwise
IfZero	$= \lambda \lambda x y z. x (\mathbf{K} z) y$	$\mathbf{IfZero} \ulcorner 0 \urcorner \ulcorner p \urcorner \ulcorner q \urcorner =_{\beta} \ulcorner p \urcorner$ $\mathbf{IfZero} \ulcorner n + 1 \urcorner \ulcorner p \urcorner \ulcorner q \urcorner =_{\beta} \ulcorner q \urcorner$

4.4 Church Lists

The Church encoding of a list xs is the term $\ulcorner xs \urcorner$ defined recursively by:

$$\begin{aligned} \ulcorner \mathbf{Nil} \urcorner &= \lambda c n. n, \\ \ulcorner x : xs \urcorner &= \lambda c n. c \ulcorner x \urcorner (\ulcorner xs \urcorner c n). \end{aligned}$$

This gives us the definition of **Cons**:

$$\mathbf{Cons} = \lambda x y. \lambda c n. c x (y c n).$$

5 Recursion

5.1 The Factorial

We have that fact is the unique function on the natural numbers that satisfies the equation:

$$F(n) = \begin{cases} 1 & \text{if } n = 0 \\ n \cdot F(n-1) & \text{otherwise.} \end{cases}$$

5.2 Fixed Points

We say that a term N is a fixed point of another term M if $MN =_{\beta} N$.

5.3 The First Recursion Theorem

Every term possesses a fixed point.

Proof. Let M be a term. We define:

$$\mathbf{Y} = \lambda f.(\lambda x.f(xx))(\lambda x.f(xx)),$$

called Curry's Paradoxical Combinator. We will show that $M(\mathbf{Y}M) =_{\beta} \mathbf{Y}M$ (so $\mathbf{Y}M$ is a fixed point of M):

$$\begin{aligned} \mathbf{Y}M &\rightarrow_{\beta} (\lambda x.M(xx))(\lambda x.M(xx)) \\ &\rightarrow_{\beta} M((\lambda x.M(xx))(\lambda x.M(xx))) \end{aligned}$$

$$M(\mathbf{Y}M) \rightarrow_{\beta} M((\lambda x.M(xx))(\lambda x.M(xx)))$$

$$M(\mathbf{Y}M) =_{\beta} \mathbf{Y}M,$$

as required. □

5.3.1 Solving Recursive Definitions

We want to solve for some term M :

$$\begin{aligned} Mx_1 \cdots x_n &=_{\beta} N[M/y] \\ \iff M &=_{\beta} \lambda x_1 \cdots x_n. N[M/y] \\ \iff M &=_{\beta} (\lambda y. \lambda x_1 \cdots x_n. N)M \end{aligned}$$

so if we find a fixed point of $(\lambda y. \lambda x_1 \cdots x_n. N)$ then we have an M which satisfies the equation. But, by the first recursion theorem, we can always find such M . Thus, we can just set $M = \mathbf{Y}(\lambda y. \lambda x_1 \cdots x_n. N)$.

6 Induction

We let Φ be some property of terms. We have that if the following conditions are met then it follows that Φ holds for all terms:

1. For all variables x , $\Phi(x)$ holds,
2. For all terms P, Q , if $\Phi(P)$ and $\Phi(Q)$ hold then $\Phi(PQ)$ holds,
3. For all terms P and variables x , if $\Phi(P)$ holds then $\Phi(\lambda x.P)$ holds.

6.1 The Induction Metaprinciple

For some set S with an inductive definition defined by rules R_1, \dots, R_k . The induction principle for proving that for all s in S , $\Phi(s)$ holds has k clauses corresponding to the rules of S .

If a rule R_i has m premises and a side condition ψ :

$$\psi \frac{s_1 \in S \cdots s_m \in S}{s \in S}, (R_i)$$

then the corresponding clause in the induction principle requires that if $\Phi(s_1), \dots, \Phi(s_m), \Phi(\psi)$ hold then $\Phi(s)$ holds.

7 Types

7.1 Monotypes

We assume a countable set of type variables \mathbb{A} (usually denoted by a, b, c , etc.). The monotypes \mathbb{T} are a set of strings defined inductively by the type variable rule:

$$\overline{a \in \mathbb{T}},$$

for some $a \in \mathbb{A}$ and the arrow rule:

$$\frac{A \in T \quad B \in T}{(A \rightarrow B) \in \mathbb{T}}.$$

7.2 Type Schemes

Type schemes are pairs consisting of a finite set of type variables a_1, \dots, a_m and a monotype A which we write as $\forall a_1 \dots a_m. A$.

7.3 Free Type Variables

We define the set of free type variables for a type scheme $\forall \bar{a}. A$, denoted by $FTV(\forall \bar{a}. A)$, recursively by the following rules:

$$\begin{aligned} FTV(a) &= \{a\}, \\ FTV(A \rightarrow B) &= FTV(A) \cup FTV(B), \\ FTV(\forall a_1 \dots a_m. A) &= FTV(A) \setminus \{a_1, \dots, a_m\}. \end{aligned}$$

We consider type schemes that only differ by choice of bound variable names to be equivalent.

7.4 Type Substitution

A type substitution is a total map σ, τ, θ from \mathbb{A} to \mathbb{T} with the property that $\sigma(a) \neq a$ for only finitely many $a \in \mathbb{A}$. We define the map as follows:

$$\begin{aligned} a\sigma &= \sigma(a) \\ (A \rightarrow B)\sigma &= A\sigma \rightarrow B\sigma. \end{aligned}$$

7.4.1 Composition of Substitutions

We write $\sigma_1\sigma_2$ for the substitution obtained by composing σ_2 after σ_1 , defined as by $(\sigma_1\sigma_2)(a) = (\sigma_1(a))\sigma_2$.

7.5 Type Assignment

A type assignment is a pair of a term M and a type scheme A written $M : A$. The term part is called the subject and the type part the predicate.

7.6 Type Environments

A type environment written Γ is a finite set of type assignments of the form $x : \forall \bar{a}. A$ which are consistent in the sense that multiple type assignments of the same subject will agree.

The subjects of Γ is the set $\text{dom}(\Gamma)$.

7.7 Type Judgement

A type judgement is a triple consisting of a type environment Γ , a term M and a monotype A written as $\Gamma \vdash M : A$.

7.8 The Type System

The type system is defined by the following rules. For $x : \forall \bar{a}. A$ in Γ we have the type variable rule:

$$\overline{\Gamma \vdash x : A[\bar{B}/\bar{a}]},$$

the rule of type application:

$$\frac{\Gamma \vdash M : B \rightarrow A \quad \Gamma \vdash N : B}{\Gamma \vdash MN : A},$$

and for x not in $\text{dom}(\Gamma)$, we have the rule of type abstraction:

$$\frac{\Gamma \cup \{x : B\} \vdash M : A}{\Gamma \vdash \lambda x. M : B \rightarrow A}.$$

A proof tree justifying a type judgement is called a type derivation.

7.9 Typability

We say a closed term M is typable if there is some type A such that $\{\} \vdash M : A$. We have that $\lambda x. xx$ is untypable.

7.10 The Inversion Theorem

Suppose $\Gamma \vdash M : A$ is derivable, we have that:

- if M is a variable x then there is a type scheme $\forall \bar{a}. B$ in Γ with $A = B[\bar{C}/\bar{a}]$ for some monotypes \bar{C} ,
- if M is an application PQ then there is a type B such that $\Gamma \vdash P : B \rightarrow A$ and $\Gamma \vdash Q : B$,
- if M is an abstraction $\lambda x. P$ then there are types B and C such that $A = B \rightarrow C$ and $\Gamma \cup \{x : B\} \vdash P : C$.

7.11 Properties of the System

For a term M , environment Γ and type S , we have the following properties.

Subterm Closure

If $\Gamma \vdash M : S$ is derivable and N is a subterm of M , there is some Γ' containing Γ and some S' such that $\Gamma' \vdash N : S'$.

Relevance (1st)

If $\Gamma \vdash M : S$ then $FV(M) \subseteq \text{dom}(\Gamma)$.

Relevance (2nd)

If $\Gamma \vdash M : S$ then $\{x : A \text{ where } x : A \in \Gamma \text{ and } x \in FV(M)\} \vdash M : S$.

Weakening

If $\Gamma \vdash M : S$ and $\Gamma \subseteq \Gamma'$ then $\Gamma' \vdash M : S$.

Subject Reduction

If $\Gamma \vdash M : A$ and $M \rightarrow_\beta N$ then $\Gamma \vdash N : A$.

Preservation under Substitution

If $\Gamma, x : B \vdash M : A$ and $\Gamma \vdash N : B$ then $\Gamma \vdash M[N/x] : A$.

Subformula

We suppose Γ has predicates solely consisting of monotypes, for some term M in β -normal form with $\Gamma \vdash M : A$ we have that the derivation of this judgement is unique and all of the types mentioned in the derivation are substrings of types mentioned in the conclusion.

7.12 Type Constraints

A type constraint is a pair of monotypes (A, B) written suggestively as $A \stackrel{?}{=} B$.

7.13 Constraint Generation

We create a constraint generating function, which takes in an environment and a term and returns a pair consisting of a type variable and a set of constraints:

$$\begin{aligned} \text{CGen}(\Gamma, x) = & \\ & \text{let } a \text{ be fresh} \\ & \text{let } \forall a_1, \dots, a_k. A = \Gamma(x) \\ & \text{let } b_1, \dots, b_k \text{ be fresh} \\ & (a, \{a \stackrel{?}{=} A[b_1/a_1, \dots, b_k/a_k]\}) \end{aligned}$$

$$\begin{aligned} \text{CGen}(\Gamma, MN) = & \\ & \text{let } a \text{ be fresh} \\ & \text{let } (b, \mathcal{C}_1) = \text{CGen}(\Gamma, M) \\ & \text{let } (c, \mathcal{C}_2) = \text{CGen}(\Gamma, N) \\ & (a, \mathcal{C}_1 \cup \mathcal{C}_2 \cup \{b \stackrel{?}{=} c \rightarrow a\}) \end{aligned}$$

$$\begin{aligned} \text{CGen}(\Gamma, \lambda x. M) = & \\ & \text{let } a \text{ be fresh} \\ & \text{let } b \text{ be fresh} \\ & \text{let } (c, \mathcal{C}) = \text{CGen}(\Gamma \cup \{x : b\}, M) \\ & (a, \mathcal{C} \cup \{a \stackrel{?}{=} b \rightarrow c\}) \end{aligned}$$

7.14 Unifier

A unifier to a finite set of type constraints $\{A_1 \stackrel{?}{=} B_1, \dots, A_n \stackrel{?}{=} B_n\}$ is a type substitution σ such that for all i in $[n]$, $A_i\sigma = B_i\sigma$.

If $\text{CGen}(\Gamma, M) = (a, \mathcal{C})$ then σ is a unifier for \mathcal{C} if and only if $\Gamma\sigma \vdash M : \sigma(a)$ is derivable.

7.14.1 Most General Unifier

For a set of type constraints \mathcal{C} , a unifier of \mathcal{C} , σ , is the most general unifier if every unifier σ^* of \mathcal{C} is of the form $\sigma\sigma'$ for some σ' .

7.15 Constraint Solved Form

A set of type constraints $\mathcal{C} = \{a_1 \stackrel{?}{=} A_1, \dots, a_m \stackrel{?}{=} A_m\}$ is in solved form if each of the a_i are distinct and do not occur in any A_j . If \mathcal{C} is in solved form, we define $[[\mathcal{C}]] = [A_1/a_1, \dots, A_m/a_m]$.

7.16 Robinson's Algorithm

Robinson's algorithm is the repeated and exhaustive application of the following rules, returning the most general unifier from a set of constraints:

$$\begin{aligned} \{A \stackrel{?}{=} A\} \uplus \mathcal{C} &\mapsto \mathcal{C}, \\ \{A_1 \rightarrow A_2 \stackrel{?}{=} B_1 \rightarrow B_2\} \uplus \mathcal{C} &\mapsto \{A_1 \stackrel{?}{=} B_1, A_2 \stackrel{?}{=} B_2\} \uplus \mathcal{C}, \\ \{A \stackrel{?}{=} a\} \uplus \mathcal{C} &\mapsto \{a \stackrel{?}{=} A\} \uplus \mathcal{C}, & (A \notin \mathbb{A}) \\ \{a \stackrel{?}{=} A\} \uplus \mathcal{C} &\mapsto \{a \stackrel{?}{=} A\} \uplus \mathcal{C}[A/a], & (a \notin FTV(A)) \end{aligned}$$

7.17 Hindley-Milner Type Inference

For an input, closed term M , we perform the following steps:

1. Generate the constraints \mathcal{C} and type variable a using $\text{CGen}(\emptyset, M)$,
2. Solve \mathcal{C} using Robinson's algorithm to obtain the most general unifier or deduce unsolvability,
3. If \mathcal{C} has no solution, M is untypable. Otherwise, we return $\sigma(a)$.

7.18 Principle Type Scheme Theorem

If a closed term M is typable, then Hindley-Milner type inference returns a type A that is principal meaning:

- $\vdash M : A$ is derivable,
- If $\vdash M : B$ is derivable, there exists monotypes C_1, \dots, C_k such that $B = A[C_1/a_1, \dots, C_k/a_k]$.